

# DRIVEN DATA: The Tanzanian Water Crisis

Author: Namita Rana

## Overview

Tanzania, as a developing country, struggles with providing clean water to its population of over 57,000,000. There are many water points already established in the country, but some are in need of repair while others have failed altogether. As a result, many people are suffering. In this project, the objective is to dive into this crisis, focusing on the analysis of water well functionality accross the country in order to assess the level of the crisis and explore some possible solutions towards improving the situation for millions of people.

## Business Problem:

Almost half the population of Tanzania is without basic access to safe water. Although there are many waterpoints already established in the country, a lot of them are in need of repair while others have failed altogether. It can assist the Tanzanian Ministry of Water on identifying pumps that are in need of repair and/or no longer functional. Understanding which pumps will fail and which will not may help improve maintenance operations and ensure that clean water is available to people residing in Tanzania. Note that this is a ternary classification problem by default, but can be engineered to be binary. Stakeholder: Ministry of Water and Irrigation, Goverment of Tanzania.

## Dataset

The dataset provided on <https://www.drivendata.org/> (<https://www.drivendata.org/>) by Taarifa and the Tanzanian Ministry of Water.

For this dataset, there are two subsections to the problem description:

### Features

- List of features

### Labels

- List of labels

### The features in this dataset are:

- amount\_tsh - Total static head (amount water available to waterpoint)
- date\_recorded - The date the row was entered

- funder - Who funded the well
- gps\_height - Altitude of the well
- installer - Organization that installed the well
- longitude - GPS coordinate
- latitude - GPS coordinate
- wpt\_name - Name of the waterpoint if there is one
- num\_private -
- basin - Geographic water basin
- subvillage - Geographic location
- region - Geographic location
- region\_code - Geographic location (coded)
- district\_code - Geographic location (coded)
- lga - Geographic location
- ward - Geographic location
- population - Population around the well
- public\_meeting - True/False
- recorded\_by - Group entering this row of data
- scheme\_management - Who operates the waterpoint
- scheme\_name - Who operates the waterpoint
- permit - If the waterpoint is permitted
- construction\_year - Year the waterpoint was constructed
- extraction\_type - The kind of extraction the waterpoint uses
- extraction\_type\_group - The kind of extraction the waterpoint uses
- extraction\_type\_class - The kind of extraction the waterpoint uses
- management - How the waterpoint is managed
- management\_group - How the waterpoint is managed
- payment - What the water costs
- payment\_type - What the water costs
- water\_quality - The quality of the water
- quality\_group - The quality of the water
- quantity - The quantity of water
- quantity\_group - The quantity of water
- source - The source of the water
- source\_type - The source of the water
- source\_class - The source of the water
- waterpoint\_type - The kind of waterpoint
- waterpoint\_type\_group - The kind of waterpoint

The labels in this dataset:

- Functional
- Non Functional
- Functional but needs repair.

## **Data Structure, Selection & Transformation:**

The raw dataset contains 60,000 water pumps and 40 different features. Combining feature selection and engineering we reduced the dataset to 14 features. We are using Functional / Non-Functional/Functional but needs repair as our target. We made some additional small features changes to simplify the dataset detailed in the Main jupyter notebook.

## **Methods**

### **Cleaning and Feature Engineering**

This project uses data cleaning and feature engineering to also addressed the class imbalance between classes we have used SMOTE. We also cleaned up our data, as there were missing values in multiple columns. For example, if "construction year" of the waterpoint had a missing value, we filled that value with the mean construction year of all waterpoints within that given region. We also engineered one column "age" to equal the age of the waterpoint when the data was collected. Cleaning our data and building these features helped make our model more interpretable and significant.

### **Models Development**

Classification was used in order to predict the status level of a pump. We trained four separate classifier models: K-Nearest Neighbor, Decision Trees, Random Forest, and XGBoost. Knn was our simplest model, and we tuned our models to get the best results.

### **Results**

Our simplest model, XGBoost came back with a confusion matrix that produced a 82% accuracy score and a 86% precision score. For our purposes, we were looking to maximize precision as we want to reduce the amount of False positives (predicting the pump as non-functional when it is functional).

## **Importing the libraries.**

```
In [1]: #import libraries
import pandas as pd
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import matplotlib.ticker as mtick
import seaborn as sns
from scipy import stats
import numpy as np
from sklearn.preprocessing import StandardScaler
from imblearn.over_sampling import SMOTE
from sklearn.impute import KNNImputer
from sklearn.metrics import precision_score, recall_score, accuracy_score
from sklearn.metrics import roc_curve, auc
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import plot_confusion_matrix
from sklearn.metrics import classification_report
from sklearn.model_selection import GridSearchCV
from sklearn.neighbors import KNeighborsClassifier
from sklearn import tree
from sklearn.ensemble import RandomForestClassifier
import xgboost as xgb
from sklearn import svm
from sklearn.metrics import f1_score, balanced_accuracy_score, plot_cc
```

## Import the data

```
In [2]: #obtain the data and read the file.
df = pd.read_csv('Tanzanian_WellDataset.csv', index_col=False)
df
```

Out[2]:

	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latit
0	69572	6000.0	2011-03-14	Roman	1390	Roman	34.938093	-9.856
1	8776	0.0	2013-03-06	Grumeti	1399	GRUMETI	34.698766	-2.147
2	34310	25.0	2013-02-25	Lottery Club	686	World vision	37.460664	-3.821
3	67743	0.0	2013-01-28	Unicef	263	UNICEF	38.486161	-11.155
4	19728	0.0	2011-07-13	Action In A	0	Artisan	31.130847	-1.825
...	...	...	...	...	...	...	...	...
59395	60739	10.0	2013-05-03	Germany Republi	1210	CES	37.169807	-3.253
59396	27263	4700.0	2011-05-07	Cefa- njombe	1212	Cefa	35.249991	-9.070
59397	37057	0.0	2011-04-11	NaN	0	NaN	34.017087	-8.750
59398	31282	0.0	2011-03-08	Malec	0	Musa	35.861315	-6.378
59399	26348	0.0	2011-03-23	World Bank	191	World	38.104048	-6.747

59400 rows × 41 columns

```
In [3]: #Selecting 20000 rows of the dataframe.
df = df.iloc[:20000,:]
df
```

Out[3]:

	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	
0	69572	6000.0	2011-03-14	Roman	1390	Roman	34.938093	.
1	8776	0.0	2013-03-06	Grumeti	1399	GRUMETI	34.698766	.
2	34310	25.0	2013-02-25	Lottery Club	686	World vision	37.460664	.
3	67743	0.0	2013-01-28	Unicef	263	UNICEF	38.486161	-1
4	19728	0.0	2011-07-13	Action In A	0	Artisan	31.130847	.
...	...	...	...	...	...	...	...	
19995	49338	100.0	2013-01-23	Concern	196	CONCERN	39.495034	-1
19996	36601	0.0	2013-01-21	Fini Water	260	FINI WATER	38.954102	.
19997	13299	0.0	2013-02-24	Netherlands	0	DWE	33.591385	.
19998	57089	0.0	2012-10-17	Government Of Tanzania	0	Government	33.103730	.
19999	61035	0.0	2013-03-03	Danida	872	DANIDA	36.047885	-1

20000 rows × 41 columns

In [4]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20000 entries, 0 to 19999
Data columns (total 41 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     20000 non-null  int64
1   amount_tsh                           20000 non-null  float64
2   date_recorded                         20000 non-null  object
3   funder                                18778 non-null  object
4   gps_height                            20000 non-null  int64
5   installer                             18768 non-null  object
6   longitude                             20000 non-null  float64
7   latitude                             20000 non-null  float64
8   wpt_name                              20000 non-null  object
9   num_private                           20000 non-null  int64
10  basin                                 20000 non-null  object
11  subvillage                            19877 non-null  object
12  region                                20000 non-null  object
13  region_code                           20000 non-null  int64
14  district_code                         20000 non-null  int64
15  lga                                    20000 non-null  object
16  ward                                  20000 non-null  object
17  population                            20000 non-null  int64
18  public_meeting                        18861 non-null  object
19  recorded_by                           20000 non-null  object
20  scheme_management                     18668 non-null  object
21  scheme_name                           10438 non-null  object
22  permit                                18969 non-null  object
23  construction_year                     20000 non-null  int64
24  extraction_type                       20000 non-null  object
25  extraction_type_group                  20000 non-null  object
26  extraction_type_class                  20000 non-null  object
27  management                            20000 non-null  object
28  management_group                       20000 non-null  object
29  payment                               20000 non-null  object
30  payment_type                           20000 non-null  object
31  water_quality                          20000 non-null  object
32  quality_group                          20000 non-null  object
33  quantity                               20000 non-null  object
34  quantity_group                         20000 non-null  object
35  source                                20000 non-null  object
36  source_type                           20000 non-null  object
37  source_class                           20000 non-null  object
38  waterpoint_type                        20000 non-null  object
39  waterpoint_type_group                  20000 non-null  object
40  status_group                           20000 non-null  object
dtypes: float64(3), int64(7), object(31)
memory usage: 6.3+ MB
```

```
In [5]: #checking for null values:  
df.isna().sum()
```

```
Out[5]: id                                0  
amount_tsh                               0  
date_recorded                            0  
funder                                   1222  
gps_height                               0  
installer                               1232  
longitude                                0  
latitude                                0  
wpt_name                                 0  
num_private                              0  
basin                                    0  
subvillage                              123  
region                                   0  
region_code                             0  
district_code                           0  
lga                                       0  
ward                                     0  
population                              0  
public_meeting                          1139  
recorded_by                             0  
scheme_management                       1332  
scheme_name                             9562  
permit                                  1031  
construction_year                       0  
extraction_type                         0  
extraction_type_group                   0  
extraction_type_class                   0  
management                             0  
management_group                       0  
payment                                 0  
payment_type                            0  
water_quality                           0  
quality_group                           0  
quantity                                0  
quantity_group                           0  
source                                  0  
source_type                             0  
source_class                            0  
waterpoint_type                         0  
waterpoint_type_group                   0  
status_group                            0  
dtype: int64
```

## Exploratory Data Analysis:

I will investigate the data, drop the columns that I will not be using and exploring the other features/columns of the dataset.



```
In [6]: #dropping the columns with the missing values:
df = df.drop(columns=['installer', 'subvillage', 'public_meeting', 'scheme'])
df
```

Out[6]:

	id	amount_tsh	date_recorded	funder	gps_height	longitude	latitude	w
0	69572	6000.0	2011-03-14	Roman	1390	34.938093	-9.856322	
1	8776	0.0	2013-03-06	Grumeti	1399	34.698766	-2.147466	
2	34310	25.0	2013-02-25	Lottery Club	686	37.460664	-3.821329	
3	67743	0.0	2013-01-28	Unicef	263	38.486161	-11.155298	N
4	19728	0.0	2011-07-13	Action In A	0	31.130847	-1.825359	
...	...	...	...	...	...	...	...	
19995	49338	100.0	2013-01-23	Concern	196	39.495034	-10.278200	Kil
19996	36601	0.0	2013-01-21	Fini Water	260	38.954102	-9.976577	
19997	13299	0.0	2013-02-24	Netherlands	0	33.591385	-3.149213	
19998	57089	0.0	2012-10-17	Government Of Tanzania	0	33.103730	-3.915889	N
19999	61035	0.0	2013-03-03	Danida	872	36.047885	-10.617697	M

20000 rows x 35 columns

```
In [7]: #dropping the columns that I won't be using:
df = df.drop(columns=['water_quality', 'quantity_group', 'extraction_type'])
df
```

Out[7]:

	id	amount_tsh	date_recorded	funder	gps_height	longitude	latitude	wp
0	69572	6000.0	2011-03-14	Roman	1390	34.938093	-9.856322	
1	8776	0.0	2013-03-06	Grumeti	1399	34.698766	-2.147466	
2	34310	25.0	2013-02-25	Lottery Club	686	37.460664	-3.821329	
3	67743	0.0	2013-01-28	Unicef	263	38.486161	-11.155298	Nz
4	19728	0.0	2011-07-13	Action In A	0	31.130847	-1.825359	
...	...	...	...	...	...	...	...	
19995	49338	100.0	2013-01-23	Concern	196	39.495034	-10.278200	Kil
19996	36601	0.0	2013-01-21	Fini Water	260	38.954102	-9.976577	
19997	13299	0.0	2013-02-24	Netherlands	0	33.591385	-3.149213	
19998	57089	0.0	2012-10-17	Government Of Tanzania	0	33.103730	-3.915889	N
19999	61035	0.0	2013-03-03	Danida	872	36.047885	-10.617697	M

20000 rows x 26 columns

```
In [8]: #Checking for null values.  
df.isna().sum()
```

```
Out[8]: id                0  
amount_tsh              0  
date_recorded           0  
funder                 1222  
gps_height              0  
longitude               0  
latitude                0  
wpt_name                0  
basin                   0  
region                  0  
region_code             0  
district_code           0  
lga                     0  
ward                    0  
population              0  
recorded_by             0  
construction_year       0  
extraction_type_class    0  
management              0  
.
```

## Exploring the columns:

### Target variable:

Let's have a look at the target.

```
In [9]: df['status_group'].value_counts()
```

```
Out[9]: functional                10878  
non functional                   7680  
functional needs repair          1442  
Name: status_group, dtype: int64
```

```
In [10]: df.status_group.value_counts(normalize=True)
```


```
Out[10]: functional                0.5439  
non functional                   0.3840  
functional needs repair          0.0721  
Name: status_group, dtype: float64
```

The master training set contains 59364 entries and 40 columns, with the target variable being `status_group`. This dataset has some class imbalances that would have to be addressed during modeling. Although this is a ternary dataset, one class "functional" takes up 54.3% of the whole training dataset, while "functional needs repair" is only 7.2%.

```
In [11]: #Plot the status_group spread
ax = df.status_group.value_counts(normalize = True).plot(kind = 'bar',
                                                         color = ['forestgreen', 'firebrick', 'yellow'],
                                                         edgecolor = 'black')

#Format x- and y-axis
plt.xticks(fontsize = 13, rotation = 0)
plt.xlabel('\nWater Pump Status', fontweight = 'bold', fontsize = 14)
ax.yaxis.set_major_formatter(mtick.PercentFormatter(1.0))
plt.yticks(fontsize = 13)
plt.ylabel('Recorded Percentage\n', fontweight = 'bold', fontsize = 14)

#Format plot
plt.title('Water Pumps Functionality Status Spread \n', fontsize = 16,
plt.savefig("status_spread.png")
```



In [12]: `df.describe()`

Out[12]:

	id	amount_tsh	gps_height	longitude	latitude	region_code
<b>count</b>	20000.00000	20000.000000	20000.00000	20000.000000	2.000000e+04	20000.000000
<b>mean</b>	37007.70645	325.394430	666.81270	34.111313	-5.715622e+00	15.480050
<b>std</b>	21580.96974	3459.998068	693.17267	6.554726	2.948522e+00	17.892049
<b>min</b>	0.00000	0.000000	-63.00000	0.000000	-1.158630e+01	1.000000
<b>25%</b>	18314.75000	0.000000	0.00000	33.103252	-8.569859e+00	5.000000
<b>50%</b>	36892.50000	0.000000	364.00000	34.912733	-5.034241e+00	12.000000
<b>75%</b>	55895.50000	20.000000	1320.00000	37.210297	-3.325097e+00	17.000000
<b>max</b>	74246.00000	350000.000000	2770.00000	40.345193	-2.000000e-08	99.000000

**Let's have a look at the amount of water in the well:  
amount\_tsh**


In [13]: *#looking at amount of water equal to zero.*  
`df[df.amount_tsh == 0].status_group.value_counts(normalize=True)`

Out[13]: functional 0.473756  
 non functional 0.453947  
 functional needs repair 0.072297  
 Name: status\_group, dtype: float64

In [14]: `len(df[df.amount_tsh == 0])`

Out[14]: 13984

```
In [15]: #Plot to see the amount of water left.
fig, ax = plt.subplots(figsize = (8, 5))
sns.stripplot(x = 'status_group', y = 'amount_tsh', data = df, edgecol=
              palette = "winter");
plt.savefig("status_group.png")
```



**Lets look at construction year:**

```
In [16]: len(df[df.construction_year == 0])/len(df)
```

```
Out[16]: 0.3478
```


```
In [17]: df[df.construction_year != 0].construction_year.describe()
```

```
Out[17]: count      13044.000000
mean        1996.787182
std          12.583458
min          1960.000000
25%          1987.000000
50%          2000.000000
75%          2008.000000
max          2013.000000
Name: construction_year, dtype: float64
```

```
In [18]: #Plot to see the construction yr of wells.
plt.figure(figsize=(7,4))
sns.distplot(df[df.construction_year != 0].construction_year);
```


/Users/Ravinder/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages/seaborn/distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```



### Let's have a look at quantity of water in the wells:


```
In [19]: ax = sns.countplot(x='quantity', hue="status_group", data=df,palette =
plt.savefig('Quantity_left.png'))
```



It can be observed that although there are enough water quantity in some wells, they are non-functional. When looking at this graph, dry quantity water points have a highly correlation with non-functionality. If the water point is dry or unknown, there is high chance thw water point is non functional. On the other hand, if the quantity is enough, there is a higher chance to find functional water points.

## let's look at the source of the wells:

```
In [20]: #Plotting to see source and their status
plt.figure(figsize=(13,8))
ax = sns.countplot(x='source', hue="status_group", data=df,palette = "
plt.savefig('Source.png')
```



## Looking at funder of the wells:



```
In [21]: #Counting the unque values.  
df['funder'].value_counts().head(10)
```


```
Out[21]: Government Of Tanzania    2975  
Danida                            1015  
Hesawa                            742  
Rwssp                             459  
World Bank                       452  
World Vision                     438  
Kkkt                             413  
Unicef                           379  
Tasaf                             310  
0                                280  
Name: funder, dtype: int64
```

This column is highly categorical column with thousands different values. So, we will take most common 10 values for future encoding.

```
In [22]: df1 = df.loc[df['funder']=='Government Of Tanzania']  
df2 = df.loc[df['funder']=='Danida']  
df3 = df.loc[df['funder']=='Hesawa']  
df4 = df.loc[df['funder']=='Rwssp']  
df5 = df.loc[df['funder']=='World']  
df6 = df.loc[df['funder']=='Kkkt']  
df7 = df.loc[df['funder']=='World Vision']  
df8 = df.loc[df['funder']=='Unicef']  
df9 = df.loc[df['funder']=='Tasaf']  
df10 = df.loc[df['funder']=='0']  
df_funder = pd.concat([df1,df2,df3,df4,df5,df6,df7,df8,df9,df10], ignore_index=True)
```

```
In [23]: #Plotting the funder of the wells.
plt.figure(figsize=(26,12))
ax = sns.countplot(x='funder', hue="status_group", data=df_funder,pale
ax.set_xticklabels(ax.get_xticklabels(),rotation=90)
```

```
Out[23]: [Text(0, 0, 'Government Of Tanzania'),
Text(1, 0, 'Danida'),
Text(2, 0, 'Hesawa'),
Text(3, 0, 'Rwssp'),
Text(4, 0, 'Kkkt'),
Text(5, 0, 'World Vision'),
Text(6, 0, 'Unicef'),
Text(7, 0, 'Tasaf'),
Text(8, 0, '0')]
```




From the above plot, we realize that most of the water points which funded by government are non-functional.

## Let's have a look at water quality in the wells:

```
In [24]: df['quality_group'].value_counts()
```

```
Out[24]: good          17040
salty          1775
unknown         650
milky          282
colored        176
fluoride        77
Name: quality_group, dtype: int64
```

```
In [25]: plt.figure(figsize=(10,5))
ax = sns.countplot(x='status_group', hue="quality_group", data=df,pale
plt.savefig("Water_Quality.png")
```




**Lets's check which wells are well maintained which take payment.**

```
In [26]: df['payment'].value_counts()
```

```
Out[26]: never pay          8492
pay per bucket      3035
pay monthly         2809
unknown             2768
pay when scheme fails 1342
pay annually        1213
other                341
Name: payment, dtype: int64
```

```
In [27]: plt.figure(figsize=(20,10))
ax = sns.countplot(x='payment', hue="status_group", data=df,palette =
plt.savefig("Payemnt.png")
```




This feature shows us what the water cost. Mostly, there are lots of non-functional water points as nobody ever paid for them.

### Which waterpoint\_type's are the most functional:

```
In [28]: df['waterpoint_type'].value_counts()
```

```
Out[28]: communal standpipe      9559
hand pump                        5926
other                           2159
communal standpipe multiple     2069
improved spring                 246
cattle trough                   37
dam                             4
Name: waterpoint_type, dtype: int64
```


```
In [29]: #plotting watertype with status_group
plt.figure(figsize=(15,8))
ax = sns.countplot(x='waterpoint_type', hue="status_group", data=df,pa
plt.savefig("Waterpoint_type.png")
```



It can be seen that waterpoint type has correlation with functionality of water points. Such that, communal standpipe has higher possibility to have functional, although communal standpipe multiple and others have higher possibility for non-functionality.


**Which management company is maintaining most functional wells.**

```
In [30]: #plotting management.
plt.figure(figsize=(15,8))
ax = sns.countplot(x='management', hue="status_group", data=df,palette
plt.savefig("management.png"))
```



**Let's have a look at region column:**

```
In [31]: plt.figure(figsize=(26,12))
ax = sns.countplot(x='region', hue="status_group", data=df,palette = "
plt.savefig("Region.png"))
```



Some regions have a higher probability of functional water well. Kilimanjaro and Arusha have Pangani basin which has a higher water point between basins. It is also seen that they have higher portions for functional wells.

```
In [32]: #taking a look at iringa it has the most functional wells.
df_iringa =df.loc[df['region']=='Iringa'] #to see the Iringa area
df_iringa
```

Out[32]:

	id	amount_tsh	date_recorded	funder	gps_height	longitude	latitude	wpt_name
0	69572	6000.0	2011-03-14	Roman	1390	34.938093	-9.856322	n
16	48451	500.0	2011-07-04	Unicef	1703	34.642439	-9.106185	Kwa Ji Mte
17	58155	0.0	2011-09-04	Unicef	1656	34.569266	-9.085515	Kwa R Ch
19	18274	500.0	2011-02-22	Danida	1763	34.508967	-9.894412	n
20	48375	200.0	2011-02-27	Twe	2216	34.473430	-9.594990	n
...	...	...	...	...	...	...	...	...
19931	56712	2000.0	2011-02-25	Anglican Church	1862	35.931047	-8.267362	n
19943	55720	500.0	2011-04-16	Roman Catholic	1853	34.747690	-9.174723	k Sain Ki
19954	64171	2500.0	2011-03-15	Shipo	1608	34.760365	-8.922258	k Sin Mg
19960	65913	100.0	2011-02-18	Unice	2277	34.059099	-9.193125	n
19983	24666	2400.0	2011-03-17	Danida	1631	34.980971	-8.591899	n

1776 rows × 26 columns

**Looking at population column:**

```
In [33]: df['population'].value_counts()
```

```
Out[33]: 0          7172
         1          2403
         200         659
         150         633
         250         580
         ...
        1831          1
        232          1
        296          1
        344          1
        663          1
        Name: population, Length: 719, dtype: int64
```

```
In [34]: df.loc[df['population']==0].groupby('status_group').count()
```

```
Out[34]:
```

	id	amount_tsh	date_recorded	funder	gps_height	longitude	latitude	wpt_id
<b>status_group</b>								
<b>functional</b>	3802	3802	3802	3406	3802	3802	3802	
<b>functional needs repair</b>	590	590	590	473	590	590	590	
<b>non functional</b>	2780	2780	2780	2592	2780	2780	2780	

3 rows × 25 columns

Some functional water points has zero population, it is weird so we will change zero population to mean



```
In [35]: # to see without zero mean and median
df.loc[df['population']!=0].describe()
```

Out[35]:

	id	amount_tsh	gps_height	longitude	latitude	region_code
<b>count</b>	12828.000000	12828.000000	12828.000000	12828.000000	12828.000000	12828.000000
<b>mean</b>	36979.809635	460.885999	965.511927	36.119379	-6.163819	16.307686
<b>std</b>	21521.490357	4287.380875	615.241504	2.571485	2.736845	21.965966
<b>min</b>	1.000000	0.000000	-63.000000	29.612507	-11.586297	2.000000
<b>25%</b>	18398.250000	0.000000	341.000000	34.728677	-8.478582	4.000000
<b>50%</b>	36899.000000	0.000000	1126.000000	36.753321	-5.860612	10.000000
<b>75%</b>	55638.250000	100.000000	1468.000000	38.007032	-3.623242	16.000000
<b>max</b>	74246.000000	350000.000000	2770.000000	40.345193	-1.094797	99.000000

```
In [36]: #replacing 0's with mean:
df['population'].replace(to_replace = 0 , value =278, inplace=True)
```

```
In [37]: df.sort_values(by='population', ascending=False).head(100).groupby('st
```

Out[37]:

	id	amount_tsh	date_recorded	funder	gps_height	longitude	latitude	wpt_name
<b>status_group</b>								
<b>functional</b>	62	62	62	60	62	62	62	
<b>functional needs repair</b>	3	3	3	2	3	3	3	
<b>non functional</b>	35	35	35	34	35	35	35	

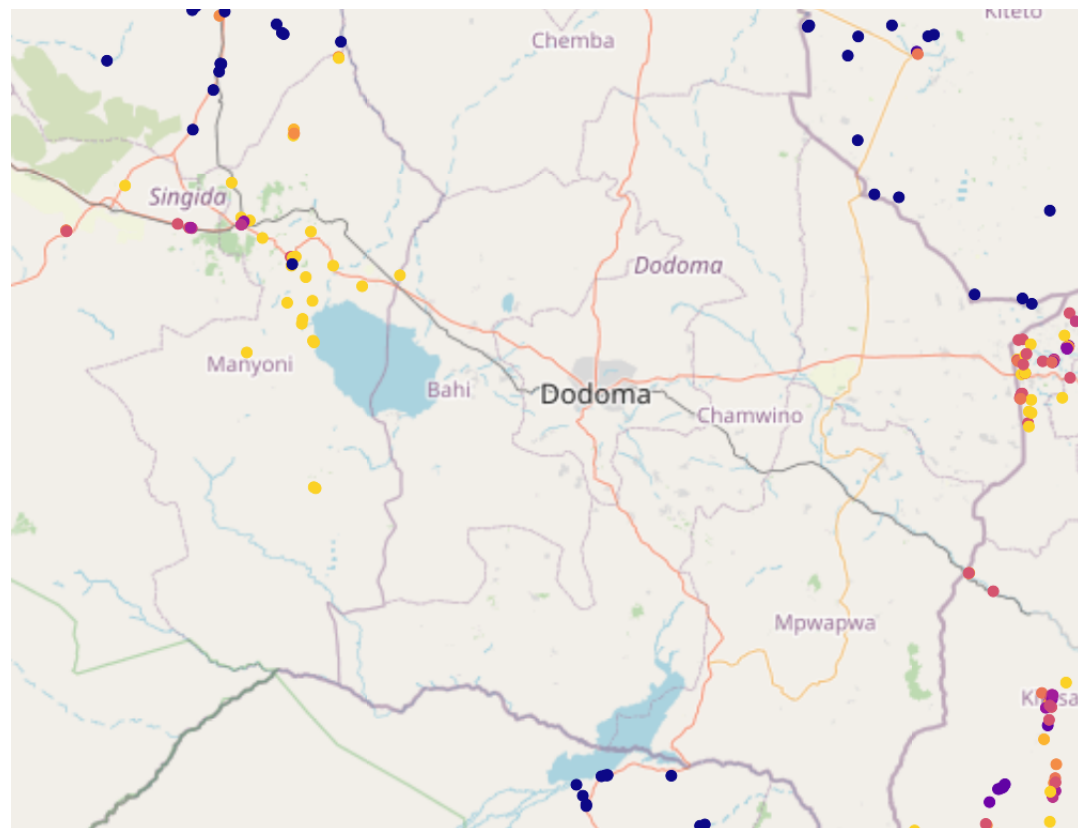
3 rows × 25 columns

```
In [38]: df['population'].mean()
```

Out[38]: 278.5583

To see the most populated areas water point functionality , we did groupby. It shows that higher population areas have more functional water points

```
In [39]: #creating a scatter plot to see where population around the well lies.  
import plotly.express as px  
fig = px.scatter_mapbox(df[df['population'] < 278],  
                        lat='latitude',  
                        lon='longitude',  
                        color='population',  
                        zoom=7)  
fig.update_layout(mapbox_style='open-street-map')  
plt.savefig('Population_around_well.pdf')  
fig.show()
```



<Figure size 432x288 with 0 Axes>

```
In [40]: #checking quality of water:
df_iringa.groupby(['quality_group', 'status_group']).count()
```

Out[40]:

		id	amount_tsh	date_recorded	funder	gps_height	longitude
quality_group status_group							
good	functional	1376	1376	1376	1374	1376	1376
	functional needs repair	42	42	42	42	42	42
	non functional	333	333	333	333	333	333
salty	functional	5	5	5	5	5	5
	non functional	4	4	4	4	4	4
unknown	functional	1	1	1	1	1	1
	non functional	15	15	15	15	15	15

When we looked at the Iringa area which has higher water points. There are also 333 wells which has soft, good water but non-functional.


### basin column:

```
In [41]: df['basin'].value_counts()
```

```
Out[41]: Lake Victoria      3449
Pangani                    3030
Rufiji                     2655
Internal                   2626
Lake Tanganyika            2152
Wami / Ruvu                2019
Lake Nyasa                 1657
Ruvuma / Southern Coast    1595
Lake Rukwa                 817
Name: basin, dtype: int64
```

### Plotting gps\_height with status\_group:

```
In [42]: fig, ax = plt.subplots(figsize=(15,10))
sns.scatterplot(data=df[df.longitude >= 25], x='latitude', y='longitude',
               size='gps_height', hue='status_group', ax=ax);
plt.savefig('HeightofWell')
```



We can see that waterpoints are more densely distributed in some regions than in others.

There seems to be a high amount of non functional wells in the Southeast and Northwest regions of Tanzania. Also, there are some large open spaces without any waterpoints being recorded. In addition, visually we can also see that more of the "larger" waterpoints (meaning they're higher in altitude gps\_height) have been recorded as functional.

## Let's have a look at categorical features:

Checking all the unique and missing values of each column:

*We will create a function and check for the unique and missing values.*

```
In [136]: #Cretaing the function.
def checking_unique_missing(list):
    '''
    check the Feature Name, Number of Unique Values, Number of Missing V
    '''
    for i in list:
        print("Feature Name:", i)
        print("Number of Unique Values:", len(df[i].unique()))
        print("Unique Values:", df[i].unique())
        print("Missing Values:", df[i].isna().sum())
        print('\n')
```

```
In [44]: checking_unique_missing(['extraction_type_class', 'payment', 'quality_gr
```

```
Feature Name: extraction_type_class
Number of Unique Values: 7
Unique Values: ['gravity' 'submersible' 'handpump' 'other' 'motorpum
p' 'wind-powered'
'rope pump']
Missing Values: 0
```

```
Feature Name: payment
Number of Unique Values: 7
Unique Values: ['pay annually' 'never pay' 'pay per bucket' 'unknown
'
'pay when scheme fails' 'other' 'pay monthly']
Missing Values: 0
```

```
Feature Name: quality_group
Number of Unique Values: 6
Unique Values: ['good' 'salty' 'milky' 'unknown' 'fluoride' 'colored
']
Missing Values: 0
```

```
Feature Name: quantity
Number of Unique Values: 5
Unique Values: ['enough' 'insufficient' 'dry' 'seasonal' 'unknown']
Missing Values: 0
```

```
Feature Name: source
Number of Unique Values: 10
Unique Values: ['spring' 'rainwater harvesting' 'dam' 'machine dbh'
'other'
'shallow well' 'river' 'hand dtw' 'lake' 'unknown']
Missing Values: 0
```

```

Feature Name: source_class
Number of Unique Values: 3
Unique Values: ['groundwater' 'surface' 'unknown']
Missing Values: 0

```

```

Feature Name: waterpoint_type
Number of Unique Values: 7
Unique Values: ['communal standpipe' 'communal standpipe multiple' '
hand pump' 'other'
'improved spring' 'cattle trough' 'dam']
Missing Values: 0

```

```

Feature Name: management
Number of Unique Values: 12
Unique Values: ['vwc' 'wug' 'other' 'private operator' 'water board'
'wua' 'company'
'water authority' 'parastatal' 'unknown' 'other - school' 'trust']
Missing Values: 0

```

```

Feature Name: region
Number of Unique Values: 21
Unique Values: ['Iringa' 'Mara' 'Manyara' 'Mtwara' 'Kagera' 'Tanga'
'Shinyanga' 'Tabora'
'Pwani' 'Ruvuma' 'Kilimanjaro' 'Rukwa' 'Mwanza' 'Kigoma' 'Lindi' 'D
odoma'
'Arusha' 'Mbeya' 'Singida' 'Morogoro' 'Dar es Salaam']
Missing Values: 0

```

```

Feature Name: basin
Number of Unique Values: 9
Unique Values: ['Lake Nyasa' 'Lake Victoria' 'Pangani' 'Ruvuma / Sou
thern Coast'
'Internal' 'Lake Tanganyika' 'Wami / Ruvu' 'Rufiji' 'Lake Rukwa']
Missing Values: 0

```

```

Feature Name: region_code
Number of Unique Values: 26
Unique Values: [11 20 21 90 18  4 17 14 60 10  3 15 19 16 80  1  6
2 12 13  5  7 99 24
 9  8]
Missing Values: 0

```

```

Feature Name: district_code
Number of Unique Values: 19
Unique Values: [ 5  2  4 63  1  8  3  6 43  7 23 33 53 62 60 30 13
0 80]
Missing Values: 0

```

In [45]: `df.describe()`

Out[45]:

	id	amount_tsh	gps_height	longitude	latitude	region_code
<b>count</b>	20000.00000	20000.000000	20000.00000	20000.000000	2.000000e+04	20000.000000
<b>mean</b>	37007.70645	325.394430	666.81270	34.111313	-5.715622e+00	15.480050
<b>std</b>	21580.96974	3459.998068	693.17267	6.554726	2.948522e+00	17.892049
<b>min</b>	0.00000	0.000000	-63.00000	0.000000	-1.158630e+01	1.000000
<b>25%</b>	18314.75000	0.000000	0.00000	33.103252	-8.569859e+00	5.000000
<b>50%</b>	36892.50000	0.000000	364.00000	34.912733	-5.034241e+00	12.000000
<b>75%</b>	55895.50000	20.000000	1320.00000	37.210297	-3.325097e+00	17.000000
<b>max</b>	74246.00000	350000.000000	2770.00000	40.345193	-2.000000e-08	99.000000

## Let's have a look at the continous features:

For the `date_recorded` column, I will split it into year, date and month.

The idea is to calculate the age of the well by: `date_recorded - construction year`

```
In [46]: #Tidy up the date-yr-month into separate columns
df[["year", "month", "day"]] = df["date_recorded"].str.split("-", expand=3)
```

Out[46]:

	id	amount_tsh	date_recorded	funder	gps_height	longitude	latitude	v
0	69572	6000.0	2011-03-14	Roman	1390	34.938093	-9.856322	
1	8776	0.0	2013-03-06	Grumeti	1399	34.698766	-2.147466	
2	34310	25.0	2013-02-25	Lottery Club	686	37.460664	-3.821329	
3	67743	0.0	2013-01-28	Unicef	263	38.486161	-11.155298	
4	19728	0.0	2011-07-13	Action In A	0	31.130847	-1.825359	
...	...	...	...	...	...	...	...	...

```
In [47]: df['construction_year'].value_counts()
```

```
Out[47]: 0          6956
2010         917
2008         893
2009         838
2000         706
2007         538
2006         504
2003         449
2011         410
2012         390
2002         366
2004         362
1978         349
2005         335
1995         318
1999         318
1998         317
1990         317
1985         296
1980         285
1984         264
1996         259
1982         257
1972         243
1994         241
1974         238
1997         227
```



```

2001    199
1992    197
1993    191
1983    173
1975    171
1988    157
1986    152
1970    146
1976    141
1991    115
1989    104
1987     97
1981     81
1977     74
1973     61
2013     59
1979     53
1971     48
1960     38
1968     35
1967     31
1963     26
1969     21
1964     15
1962     13
1961      4
1965      3
1966      2

```

Name: construction\_year, dtype: int64

In [48]: df.describe()

Out[48]:

	id	amount_tsh	gps_height	longitude	latitude	region_code
<b>count</b>	20000.00000	20000.000000	20000.00000	20000.000000	2.000000e+04	20000.000000
<b>mean</b>	37007.70645	325.394430	666.81270	34.111313	-5.715622e+00	15.480050
<b>std</b>	21580.96974	3459.998068	693.17267	6.554726	2.948522e+00	17.892049
<b>min</b>	0.00000	0.000000	-63.00000	0.000000	-1.158630e+01	1.000000
<b>25%</b>	18314.75000	0.000000	0.00000	33.103252	-8.569859e+00	5.000000
<b>50%</b>	36892.50000	0.000000	364.00000	34.912733	-5.034241e+00	12.000000
<b>75%</b>	55895.50000	20.000000	1320.00000	37.210297	-3.325097e+00	17.000000
<b>max</b>	74246.00000	350000.000000	2770.00000	40.345193	-2.000000e-08	99.000000

Since, construction\_year has lot of 0 values, to find the approx age of the wells, I am filling the 0 values with the median of the value: 1986, we will assume these wells are constructed after 1986.

```
In [49]: df['construction_year']=df['construction_year'].replace(0,1986)
df
```

Out[49]:

	id	amount_tsh	date_recorded	funder	gps_height	longitude	latitude	w
0	69572	6000.0	2011-03-14	Roman	1390	34.938093	-9.856322	
1	8776	0.0	2013-03-06	Grumeti	1399	34.698766	-2.147466	
2	34310	25.0	2013-02-25	Lottery Club	686	37.460664	-3.821329	
3	67743	0.0	2013-01-28	Unicef	263	38.486161	-11.155298	N
4	19728	0.0	2011-07-13	Action In A	0	31.130847	-1.825359	
...	...	...	...	...	...	...	...	
19995	49338	100.0	2013-01-23	Concern	196	39.495034	-10.278200	Kil
19996	36601	0.0	2013-01-21	Fini Water	260	38.954102	-9.976577	
19997	13299	0.0	2013-02-24	Netherlands	0	33.591385	-3.149213	
19998	57089	0.0	2012-10-17	Government Of Tanzania	0	33.103730	-3.915889	N
19999	61035	0.0	2013-03-03	Danida	872	36.047885	-10.617697	M

20000 rows x 29 columns

```
In [50]: #Finding the age of well:
df["Age_of_Well"] = df.apply(lambda x: x['year'] - x['construction_year'], axis=1)
df
```

Out[50]:

	id	amount_tsh	date_recorded	funder	gps_height	longitude	latitude	v
0	69572	6000.0	2011-03-14	Roman	1390	34.938093	-9.856322	
1	8776	0.0	2013-03-06	Grumeti	1399	34.698766	-2.147466	
2	34310	25.0	2013-02-25	Lottery Club	686	37.460664	-3.821329	
3	67743	0.0	2013-01-28	Unicef	263	38.486161	-11.155298	
4	19728	0.0	2011-07-13	Action In A	0	31.130847	-1.825359	
...	...	...	...	...	...	...	...	...

```
In [51]: #Viewing construction_yr of wells and binning them:
view_age = [df.construction_year > 2005,
             df.construction_year > 2000,
             df.construction_year > 1990,
             df.construction_year > 1980,
             df.construction_year > 1970]
vals = ['after05', '00s', '90s', '80s', '70s']


df['year_built'] = np.select(view_age, vals, 'others')
```

```
In [52]: #Plot to view the condition of wells according to their construction_y
fig = plt.figure(figsize = (15,8))

sns.countplot(x = 'year_built', hue = 'status_group', data = df,
              order = ['others', '70s', '80s', '90s', '00s', 'after05'])

plt.xticks(rotation = 45, ha = 'right', ticks = range(0, 6),
           labels = ['others', '1970s', '1980s', '1990s', '2000-2005',
                    'after2005'])
plt.xlabel('Built Year')
plt.ylabel('counts')
plt.legend(bbox_to_anchor = [1, 1])
fig.patch.set_visible(False)
plt.savefig("Built_year")

plt.show()
```



It seems condition of wells built before 1980's is deteriorating and are mostly non-functional.

```
id
  73284      0.00005
53999      0.00005
8929       0.00005
51079      0.00005
35556      0.00005
Name: id, dtype: float64
```

```
amount_tsh
0.0      0.69920
500.0    0.05405
50.0     0.04320
20.0     0.02585
1000.0   0.02500
Name: amount_tsh, dtype: float64
```

```
date_recorded      ^ ^ ^ ^ ^ ^ ^ ^
```

- population - 0,1
- construction\_year - 0
- payment - unknown
- quality\_group - unknown
- quantity - unknown
- extraction\_type\_class - other
- source\_class - unknown
- waterpoint\_type - other

I will separate continous features and categorical features in order to perform dummy encoding.

```
In [54]: #choosing columns for our dataset.  
df.columns
```

```
Out[54]: Index(['id', 'amount_tsh', 'date_recorded', 'funder', 'gps_height',  
               'longitude', 'latitude', 'wpt_name', 'basin', 'region', 'region_code',  
               'district_code', 'lga', 'ward', 'population', 'recorded_by',  
               'construction_year', 'extraction_type_class', 'management', 'payment',  
               'quality_group', 'quantity', 'source', 'source_class',  
               'waterpoint_type', 'status_group', 'year', 'month', 'day',  
               'Age_of_Well', 'year_built'],  
              dtype='object')
```

Dropping the columns that have too many unique values, like region code, district code.

```
In [55]: #Dropping the columns that don't need.
df_new = df.drop(['status_group', 'id', 'date_recorded', 'longitude', 'lat
df_new
```

Out[55]:

	amount_tsh	funder	gps_height	basin	region	population	recorded_by	cc
0	6000.0	Roman	1390	Lake Nyasa	Iringa	109	GeoData Consultants Ltd	
1	0.0	Grumeti	1399	Lake Victoria	Mara	280	GeoData Consultants Ltd	
2	25.0	Lottery Club	686	Pangani	Manyara	250	GeoData Consultants Ltd	
3	0.0	Unicef	263	Ruvuma / Southern Coast	Mtwara	58	GeoData Consultants Ltd	
4	0.0	Action In A	0	Lake Victoria	Kagera	278	GeoData Consultants Ltd	
...	...	...	...	...	...	...	...	...
19995	100.0	Concern	196	Ruvuma / Southern Coast	Lindi	1885	GeoData Consultants Ltd	
19996	0.0	Fini Water	260	Ruvuma / Southern Coast	Lindi	650	GeoData Consultants Ltd	
19997	0.0	Netherlands	0	Lake Victoria	Shinyanga	278	GeoData Consultants Ltd	
19998	0.0	Government Of Tanzania	0	Internal	Shinyanga	278	GeoData Consultants Ltd	
19999	0.0	Danida	872	Ruvuma / Southern Coast	Ruvuma	180	GeoData Consultants Ltd	

20000 rows × 19 columns

```
In [56]: #Define target
target = df['status_group']
```

```
In [57]: df_new.columns
```

```
Out[57]: Index(['amount_tsh', 'funder', 'gps_height', 'basin', 'region', 'population',  
               'recorded_by', 'construction_year', 'extraction_type_class',  
               'management', 'payment', 'quality_group', 'quantity', 'source',  
               'source_class', 'waterpoint_type', 'year', 'Age_of_Well', 'year_built'],  
              dtype='object')
```

```
In [58]: #Defining categorical and continuous features  
continous_feats = ['amount_tsh', 'gps_height',  
                  'population', 'construction_year']  
categorical_feats = list(df_new.drop(continous_feats, axis = 1).columns)
```

```
In [59]: categorical_feats
```

```
Out[59]: ['funder',  
          'basin',  
          'region',  
          'recorded_by',  
          'extraction_type_class',  
          'management',  
          'payment',  
          'quality_group',  
          'quantity',  
          'source',  
          'source_class',  
          'waterpoint_type',  
          'year',  
          'Age_of_Well',  
          'year_built']
```

```
In [60]: #Creating dummy variables:  
dummy = pd.get_dummies(df[categorical_feats], drop_first = True)  
dummy.shape
```

```
Out[60]: (20000, 1224)
```



```
In [61]: #Finalise df_new for modeling
df_new.drop(categorical_feats, axis = 1, inplace = True)
#df_new.drop(to_remove, axis = 1, inplace = True)

df_new = pd.concat ([df_new, dummy], axis = 1)
df_new.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20000 entries, 0 to 19999
Columns: 1228 entries, amount_tsh to year_built_others
dtypes: float64(1), int64(5), uint8(1222)
memory usage: 24.2 MB
```

```
In [62]: #Let's view this new dataframe
df_new.head(5)
```

Out[62]:

	amount_tsh	gps_height	population	construction_year	year	Age_of_Well	funder_A/co Germany	fu
0	6000.0	1390	109	1999	2011	12	0	
1	0.0	1399	280	2010	2013	3	0	
2	25.0	686	250	2009	2013	4	0	
3	0.0	263	58	1986	2013	27	0	
4	0.0	0	278	1986	2011	25	0	

5 rows × 1228 columns

## Modelling the data:

```
In [63]: #Splitting the data into test,train:
X_train, X_test, y_train, y_test = train_test_split(df_new, target, test_size=0.2)
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

Out[63]: ((15000, 1228), (5000, 1228), (15000,), (5000,))

```
In [64]: #Instantial StandardScaler
scaler = StandardScaler()

#Transform the train and test sets
X_train = scaler.fit_transform(X_train)
X_test = scaler.fit_transform(X_test)

#Convert to DataFrae
X_train = pd.DataFrame(X_train, columns = df_new.columns)
X_train.head()
```

Out[64]:

	amount_tsh	gps_height	population	construction_year	year	Age_of_Well	funder_A/c German
0	-0.034967	0.843902	-0.626755	0.869737	1.119035	-0.784054	-0.020000
1	-0.085982	1.013189	0.047396	1.483314	1.119035	-1.404733	-0.020000
2	-0.085982	1.237458	-0.563624	1.220353	1.119035	-1.138728	-0.020000
3	0.296630	-0.315063	-0.626755	0.606775	1.119035	-0.518048	-0.020000
4	-0.034967	-0.326638	-0.178073	-0.708033	-0.953508	0.634642	-0.020000

5 rows × 1228 columns

## Model Building:

***For this project, I will be building several models using different classifiers and then compare the performance metrics to choose the best classifier.***

The classifier's the data will be tested on:

- K-Nearest Neighbour
- Decision Tree Classifier
- Random Forest Classifier
- eXtreme Gradient Boosting (XGBoost)

## Evaluation Metrics:

- Precision
- Recall
- Accuracy
- F1 Score
- We will also be looking at the confusion matrix.

Creating function to compare scores of all the models:

```
In [65]: models_scores = pd.DataFrame(columns = ['Model', 'Precision', 'Recall']

# Creating the function
def calculate_scores(model, y_test, test_pred):
    '''
    parameters : (model:classifer model, y true value, y predicted value)

    Save precision, recall, accuracy, and f1 scores to a dataframe
    '''
    global scores_df

    # Calculating test scores
    precision = precision_score(y_test, test_pred, average='weighted')
    recall = recall_score(y_test, test_pred, average='weighted')
    accuracy = accuracy_score(y_test, test_pred)
    f1 = f1_score(y_test, test_pred, average='weighted')

    # Adding the scores into the dataframe
    models_scores.append({'Model':model, 'Precision':precision, 'Recall':
```

## Check for Class Imbalance Issue:

```
In [66]: print(y_train.value_counts(normalize=True))
```

```
functional          0.5420
non functional      0.3848
functional needs repair  0.0732
Name: status_group, dtype: float64
```

The functional category makes up 54% and the non functional category makes up 38% of the training dataset. The problem lies with the functional needs repair category which make up only 7% of the dataset. Since it is important to address this imbalance we will try to fix this imbalance by using the SMOTE(Synthetic Minority Oversampling) tool.

## Using SMOTE

```
In [67]: #Instantiate and train
smote = SMOTE()

X_smote, y_smote = smote.fit_sample(X_train, y_train)
```

```
In [68]: print(y_smote.value_counts(normalize=True))
```

```
functional          0.333333
non functional      0.333333
functional needs repair 0.333333
Name: status_group, dtype: float64
```

Now, the all the categories have the same distribution of 33.33% each.

```
In [69]: #new test, train, split with smote
X_train, X_test, y_train, y_test = train_test_split(X_smote, y_smote,
```

## We are starting with KNN as our baseline model.

```
In [70]: #Instantiate Classifier
KNN = KNeighborsClassifier()

#Fit model
KNN.fit(X_train, y_train)
```

```
Out[70]: KNeighborsClassifier()
```

```
In [71]: #Predict
train_pred = KNN.predict(X_train)
test_pred = KNN.predict(X_test)
```

```
In [135]: import sklearn.metrics as metrics
from sklearn.metrics import accuracy_score

print(metrics.accuracy_score(y_test, test_pred))
#Print Classification Report
print('Training Data')
print(classification_report(y_train, train_pred))
print('\n')
print('Testing Data')
print(classification_report(y_test, test_pred))
# Update the scores dataframe
calculate_scores('KNN', y_test, test_pred)
```

0.8261725155788783

Training Data

	precision	recall	f1-score	support
functional	0.81	0.88	0.84	6160
functional needs repair	0.89	0.93	0.91	6076
non functional	0.91	0.80	0.85	6056
accuracy			0.87	18292
macro avg	0.87	0.87	0.87	18292
weighted avg	0.87	0.87	0.87	18292

Testing Data


	precision	recall	f1-score	support
functional	0.76	0.83	0.79	1970
functional needs repair	0.86	0.91	0.88	2054
non functional	0.87	0.74	0.80	2074
accuracy			0.83	6098
macro avg	0.83	0.83	0.83	6098
weighted avg	0.83	0.83	0.83	6098

```
In [73]: #Plot a confusion matrix
from sklearn.metrics import confusion_matrix
from sklearn.metrics import plot_confusion_matrix
#plot_confusion_matrix(KNN, X_test, y_test)

disp = plot_confusion_matrix(KNN, X_test, y_test,
                             cmap=plt.cm.Blues,
                             xticks_rotation='vertical')
disp.ax_.set_title('K-Nearest Neighbors Classifier')

print(disp.confusion_matrix)
```

```
[[1247  357  366]
 [ 157 1836   61]
 [ 372  239 1463]]
```



Comments: The K-Nearest Neighbors classifier doesn't seem to perform that well in all the three categories. The performance metrics for the test data is worse than it is for the training data, especially for the 'functional needs repair' category. This perhaps hints at overfitting. If the f-1 score for the 'functional needs repairs' category is low for all models, I would need to pay more attention to the recall score of this category which I believe is a little more important than the precision score is for our purpose.

After running the baseline model I will be hypertuning models using GridSearchCv to see how they perform.

**Now we will run K-Nearest Neighbors using GridSearchCV.**

```
In [74]: from sklearn.model_selection import GridSearchCV
```

```
In [76]: param_grid = {
          'n_neighbors': [1,5,10], # default 5
          'weights': ['uniform','distance'], #default uniform
          'metric': ['eculidean','manhattan']
        }
knn = KNeighborsClassifier()
grid_search = GridSearchCV(knn, param_grid, verbose =1, cv = 3, scoring
grid_search.fit(X_train,y_train)
```

Fitting 3 folds for each of 12 candidates, totalling 36 fits

[Parallel(n\_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.

[Parallel(n\_jobs=-1)]: Done 36 out of 36 | elapsed: 11.4min finished

```
Out[76]: GridSearchCV(cv=3, estimator=KNeighborsClassifier(), n_jobs=-1,
                    param_grid={'metric': ['eculidean', 'manhattan'],
                                'n_neighbors': [1, 5, 10],
                                'weights': ['uniform', 'distance']}},
                    scoring='accuracy', verbose=1)
```

```
In [ ]: grid_search.best_params_
```

```
In [77]: #tuned knn
knn_tuned = KNeighborsClassifier(n_neighbors = 10, weights = 'distance',
knn_tuned.fit(X_train,y_train)
```

```
Out[77]: KNeighborsClassifier(leaf_size=10, n_neighbors=10, weights='distance
')
```

```
In [78]: #predicting
test_pred = knn_tuned.predict(X_test)
train_pred = knn_tuned.predict(X_train)
```

```
In [133]: import sklearn.metrics as metrics
from sklearn.metrics import accuracy_score
print(metrics.accuracy_score(y_test, test_pred))
#Print Classification Report
print('Training Data')
print(classification_report(y_train, train_pred))
print('\n')
print('Testing Data')
print(classification_report(y_test, test_pred))
# Update the scores dataframe
calculate_scores('K-Nearest Neighbor tuned', y_test, test_pred)
```

0.8261725155788783

Training Data

	precision	recall	f1-score	support
functional	0.81	0.88	0.84	6160
functional needs repair	0.89	0.93	0.91	6076
non functional	0.91	0.80	0.85	6056
accuracy			0.87	18292
macro avg	0.87	0.87	0.87	18292
weighted avg	0.87	0.87	0.87	18292

Testing Data

	precision	recall	f1-score	support
functional	0.76	0.83	0.79	1970
functional needs repair	0.86	0.91	0.88	2054
non functional	0.87	0.74	0.80	2074
accuracy			0.83	6098
macro avg	0.83	0.83	0.83	6098
weighted avg	0.83	0.83	0.83	6098

```
-----
-----
NameError                                Traceback (most recent call last)
<ipython-input-133-5a779c78b499> in <module>
      10 # Update the scores dataframe
      11 calculate_scores('K-Nearest Neighbor tuned', y_test,
test_pred)
----> 12 model_scores

NameError: name 'model_scores' is not defined
```

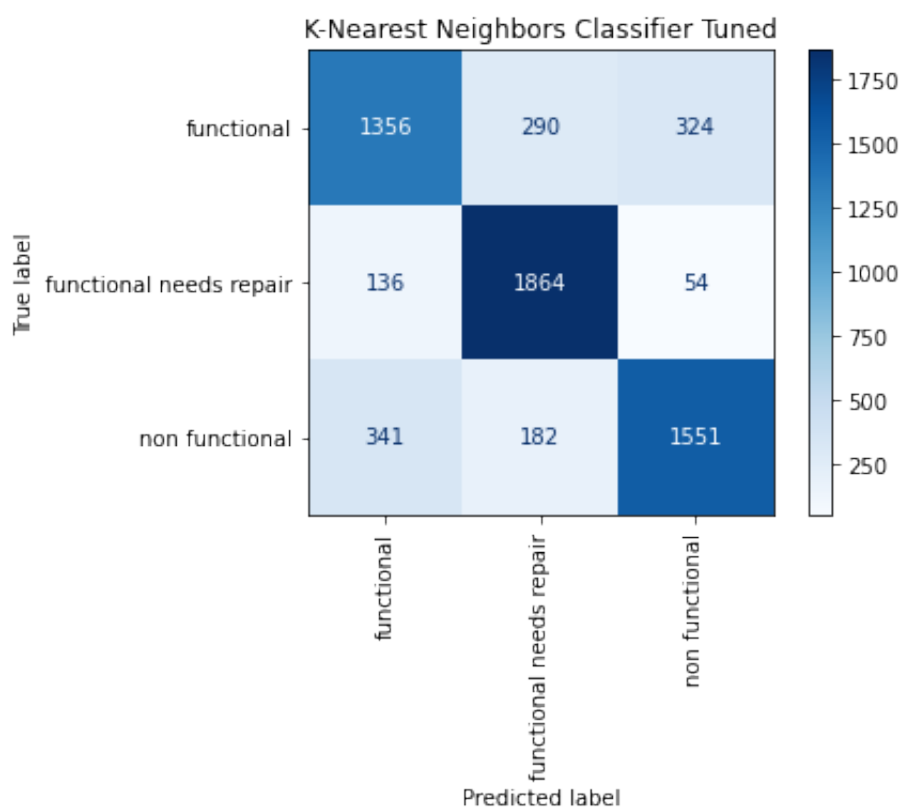


```
In [81]: #Plot a confusion matrix
from sklearn.metrics import confusion_matrix
from sklearn.metrics import plot_confusion_matrix
#plot_confusion_matrix(KNN, X_test, y_test)

disp = plot_confusion_matrix(knn_tuned, X_test, y_test,
                             cmap=plt.cm.Blues,
                             xticks_rotation='vertical')
disp.ax_.set_title('K-Nearest Neighbors Classifier Tuned')

print(disp.confusion_matrix)
```

```
[[1356  290  324]
 [ 136 1864   54]
 [ 341  182 1551]]
```



Comments: Knn tuned is performing better than the base one in terms to accuracy as well as precision. The no of false positives is less in knn\_tuned, for functional needs repair and non functionals are less.

## Decision Tree Classifier

The next model we will be taking a look is a single Decision Tree Classifier. I will be performing a grid search for criterion, max\_depth, min\_samples\_split and also min\_samples\_leaf.

```
In [82]: from sklearn.model_selection import train_test_split, GridSearchCV, cr
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.metrics import accuracy_score
```

```
In [83]: X_train, X_test, y_train, y_test = train_test_split(X_smote, y_smote,
```

```
In [84]: # Create the classifier, fit it on the training data and make predictions
decision_treeclf = DecisionTreeClassifier(criterion='entropy')

decision_treeclf.fit(X_train, y_train)
```

```
Out[84]: DecisionTreeClassifier(criterion='entropy')
```

```
In [85]: test_pred = decision_treeclf.predict(X_test)
train_pred = decision_treeclf.predict(X_train)
```

```
In [86]: print(metrics.accuracy_score(y_test, test_pred))
#Print Classification Report
print('Training Data')
print(classification_report(y_train, train_pred))
print('\n')
print('Testing Data')
print(classification_report(y_test, test_pred))
# Update the scores dataframe
calculate_scores('Decision Tree', y_test, test_pred)
```

0.7940308297802559

Training Data

	precision	recall	f1-score	support
functional	0.95	0.97	0.96	6160
functional needs repair	0.96	0.98	0.97	6076
non functional	0.99	0.95	0.97	6056
accuracy			0.97	18292
macro avg	0.97	0.97	0.97	18292
weighted avg	0.97	0.97	0.97	18292

Testing Data

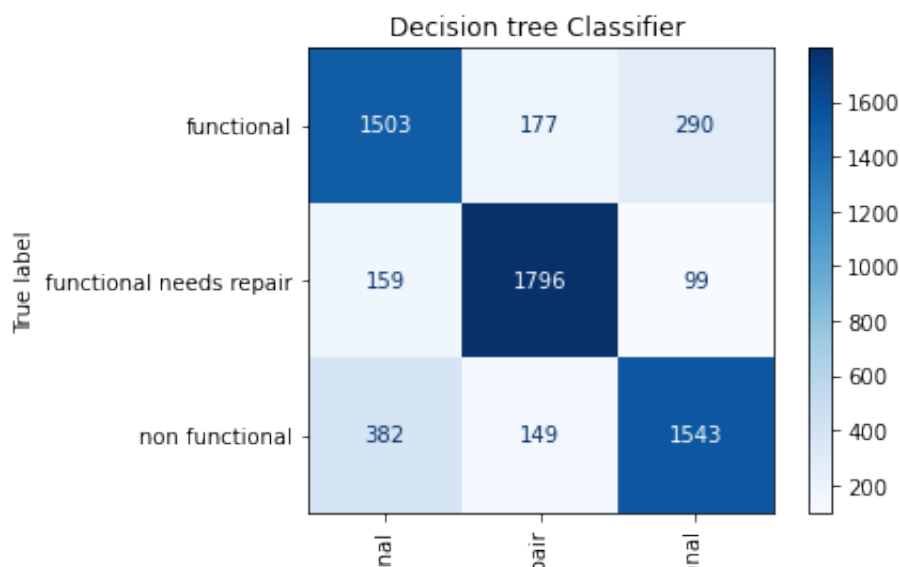
	precision	recall	f1-score	support
functional	0.74	0.76	0.75	1970
functional needs repair	0.85	0.87	0.86	2054
non functional	0.80	0.74	0.77	2074
accuracy			0.79	6098
macro avg	0.79	0.79	0.79	6098
weighted avg	0.79	0.79	0.79	6098

```
In [87]: #Plot a confusion matrix
from sklearn.metrics import confusion_matrix
from sklearn.metrics import plot_confusion_matrix
#plot_confusion_matrix(KNN, X_test, y_test)

disp = plot_confusion_matrix(decision_treeclf, X_test, y_test,
                             cmap=plt.cm.Blues,
                             xticks_rotation='vertical')
disp.ax_.set_title('Decision tree Classifier ')

print(disp.confusion_matrix)
```

```
[[1503  177  290]
 [ 159 1796   99]
 [ 382  149 1543]]
```



Decision Tree has performed better than Knn\_Tuned in terms of accuracy but in terms of precision and no of false positives, there doesn't seems much improvement. So, we will be tuning it using GridsearchCV to see if those numbers change.

### Hyperparameter Tuning of Decision Tree Classifier Using GridSearchCV:

```
In [89]: #Hypertuning Decision_Tree
dt_clf = DecisionTreeClassifier()

dt_cv_score = cross_val_score(dt_clf, X_train, y_train, cv=3)
mean_dt_cv_score = np.mean(dt_cv_score)

print(f"Mean Cross Validation Score: {mean_dt_cv_score :.2%}")
```

Mean Cross Validation Score: 76.96%

```
In [90]: dt_param_grid = {
    'criterion': ['gini', 'entropy'],
    'max_depth': [None, 2, 3, 4, 5, 6],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 3, 4, 5, 6]
}
# Instantiate GridSearchCV
dt_grid_search = GridSearchCV(dt_clf, dt_param_grid, cv=3, return_train_score=True)

# Fit to the data
dt_grid_search.fit(X_train, y_train)
```

```
Out[90]: GridSearchCV(cv=3, estimator=DecisionTreeClassifier(),
    param_grid={'criterion': ['gini', 'entropy'],
    'max_depth': [None, 2, 3, 4, 5, 6],
    'min_samples_leaf': [1, 2, 3, 4, 5, 6],
    'min_samples_split': [2, 5, 10]},
    return_train_score=True)
```

```
In [91]: # Mean training score
dt_gs_training_score = np.mean(dt_grid_search.cv_results_['mean_train_score'])

# Mean test score
dt_gs_testing_score = dt_grid_search.score(X_test, y_test)

print(f"Mean Training Score: {dt_gs_training_score :.2%}")
print(f"Mean Test Score: {dt_gs_testing_score :.2%}")
print("Best Parameter Combination Found During Grid Search:")
dt_grid_search.best_params_
```

Mean Training Score: 60.86%

Mean Test Score: 79.27%

Best Parameter Combination Found During Grid Search:

```
Out[91]: {'criterion': 'gini',
    'max_depth': None,
    'min_samples_leaf': 1,
    'min_samples_split': 2}
```

```
In [92]: test_pred = dt_grid_search.predict(X_test)
    train_pred = dt_grid_search.predict(X_train)
```

```
In [93]: print(metrics.accuracy_score(y_test, test_pred))
#Print Classification Report
print('Training Data')
print(classification_report(y_train, train_pred))
print('\n')
print('Testing Data')
print(classification_report(y_test, test_pred))
# Update the scores dataframe
calculate_scores('Decision Tree tuned', y_test, test_pred)
```

0.7927189242374549

Training Data

	precision	recall	f1-score	support
functional	0.95	0.97	0.96	6160
functional needs repair	0.96	0.98	0.97	6076
non functional	0.99	0.95	0.97	6056
accuracy			0.97	18292
macro avg	0.97	0.97	0.97	18292
weighted avg	0.97	0.97	0.97	18292

Testing Data


	precision	recall	f1-score	support
functional	0.74	0.76	0.75	1970
functional needs repair	0.84	0.87	0.86	2054
non functional	0.79	0.75	0.77	2074
accuracy			0.79	6098
macro avg	0.79	0.79	0.79	6098
weighted avg	0.79	0.79	0.79	6098

```
In [95]: #Plot a confusion matrix
from sklearn.metrics import confusion_matrix
from sklearn.metrics import plot_confusion_matrix
#plot_confusion_matrix(KNN, X_test, y_test)

disp = plot_confusion_matrix(dt_grid_search, X_test, y_test,
                             cmap=plt.cm.Blues,
                             xticks_rotation='vertical')
disp.ax_.set_title('Decision tree Classifier Tuned')

print(disp.confusion_matrix)
```

```
[[1501  166  303]
 [ 157 1783  114]
 [ 361  163 1550]]
```



Decision Tree classifier\_Tuned accuracy score is same as the base Decision Tree, doesn't show much improvement.

We will try another classifier.

## Random Forest Classifier:

```
In [96]: #Instantiate RandomForest Classifier
forest = RandomForestClassifier()
```

```
In [98]: #Use GridSerchCV to determine best parameters
rf_param_grid = {'n_estimators': [100, 200, 500],
                  'criterion': ['gini', 'entropy'],
                  'max_depth': [3, 4, 5, 6],
                  'min_samples_split': [2, 5, 10],
                  'min_samples_leaf': [2, 5] }

#Instantiate GridSearchCV()
rf_grid_search = GridSearchCV(forest, rf_param_grid, cv = 3,n_jobs = -1)

#Fit to the data
rf_grid_search.fit(X_train, y_train)
```

```
Out[98]: GridSearchCV(cv=3, estimator=RandomForestClassifier(), n_jobs=-1,
                    param_grid={'criterion': ['gini', 'entropy'],
                                'max_depth': [3, 4, 5, 6], 'min_samples_leaf': [2, 5],
                                'min_samples_split': [2, 5, 10],
                                'n_estimators': [100, 200, 500]})
```

```
In [342]: rf_grid_search.best_params_
```

```
Out[342]: {'criterion': 'gini',
           'max_depth': 6,
           'min_samples_leaf': 5,
           'min_samples_split': 2,
           'n_estimators': 200}
```

```
In [99]: rf_tuned = RandomForestClassifier(n_estimators = 500,max_depth = 6,min
rf_tuned.fit(X_train,y_train)
```

```
Out[99]: RandomForestClassifier(max_depth=6, min_samples_leaf=2, min_samples_split=5,
                                n_estimators=500)
```

```
In [100]: test_pred = rf_tuned.predict(X_test)
train_pred = rf_tuned.predict(X_train)
```



```
In [101]: print(metrics.accuracy_score(y_test, test_pred))
#Print Classification Report
print('Training Data')
print(classification_report(y_train, train_pred))
print('\n')
print('Testing Data')
print(classification_report(y_test, test_pred))
# Update the scores dataframe
calculate_scores('Random forest Classifier', y_test, test_pred)
```

0.6564447359790095

Training Data

	precision	recall	f1-score	support
functional	0.67	0.67	0.67	6160
functional needs repair	0.65	0.79	0.71	6076
non functional	0.70	0.55	0.61	6056
accuracy			0.67	18292
macro avg	0.67	0.67	0.66	18292
weighted avg	0.67	0.67	0.66	18292

Testing Data


	precision	recall	f1-score	support
functional	0.64	0.66	0.65	1970
functional needs repair	0.65	0.79	0.71	2054
non functional	0.69	0.52	0.59	2074
accuracy			0.66	6098
macro avg	0.66	0.66	0.65	6098
weighted avg	0.66	0.66	0.65	6098

```
In [346]: #Plot a confusion matrix
from sklearn.metrics import confusion_matrix
from sklearn.metrics import plot_confusion_matrix
#plot_confusion_matrix(KNN, X_test, y_test)

disp = plot_confusion_matrix(rf_tuned, X_test, y_test,
                             cmap=plt.cm.Blues,
                             xticks_rotation='vertical')
disp.ax_.set_title('Random Classifier Tuned')

print(disp.confusion_matrix)
```

```
[[1331  370  269]
 [ 357 1496  201]
 [ 520  333 1221]]
```



Random forest is not performing well for accuracy nor for the precision.

**We will try XG BOOST:**

```
In [105]: from xgboost import XGBClassifier
```

```
In [107]: X_train, X_test, y_train, y_test = train_test_split(X_smote, y_smote,
```

```
In [114]:
```

```
pile(r"\[|\]|<", re.IGNORECASE)

s = [regex.sub("_", col) if any(x in str(col) for x in set([' ', ']', ' ',
```

```
In [115]: baselinexgb = XGBClassifier()
baselinexgb.fit(X_train,y_train)
```

```
Out[115]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                        colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_i
d=-1,
                        importance_type='gain', interaction_constraints='',
                        learning_rate=0.300000012, max_delta_step=0, max_depth
=6,
                        min_child_weight=1, missing=nan, monotone_constraints=
'() ',
                        n_estimators=100, n_jobs=0, num_parallel_tree=1,
                        objective='multi:softprob', random_state=0, reg_alpha=
0,
                        reg_lambda=1, scale_pos_weight=None, subsample=1,
                        tree_method='exact', validate_parameters=1, verbosity=
None)
```

```
In [120]: import re
```

```
regex = re.compile(r"\[|\]|<", re.IGNORECASE)

X_test.columns = [regex.sub("_", col) if any(x in str(col) for x in se

train_pred = baselinexgb.predict(X_train)
test_pred = baselinexgb.predict(X_test)
```

```
In [122]: print(metrics.accuracy_score(y_test, test_pred))
#Print Classification Report
print('Training Data')
print(classification_report(y_train, train_pred))
print('\n')
print('Testing Data')
print(classification_report(y_test, test_pred))
# Update the scores dataframe
calculate_scores('XGBClassifier', y_test, test_pred)
```

0.8261725155788783

Training Data

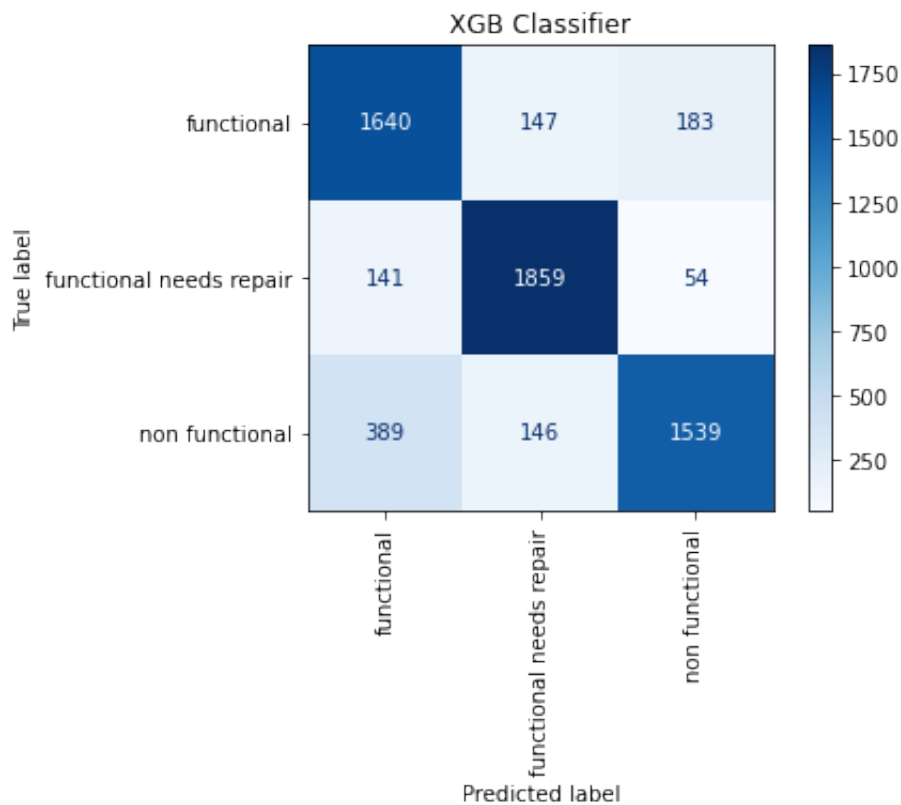
	precision	recall	f1-score	support
functional	0.81	0.88	0.84	6160
functional needs repair	0.89	0.93	0.91	6076
non functional	0.91	0.80	0.85	6056
accuracy			0.87	18292
macro avg	0.87	0.87	0.87	18292
weighted avg	0.87	0.87	0.87	18292

Testing Data

	precision	recall	f1-score	support
functional	0.76	0.83	0.79	1970
functional needs repair	0.86	0.91	0.88	2054
non functional	0.87	0.74	0.80	2074
accuracy			0.83	6098
macro avg	0.83	0.83	0.83	6098
weighted avg	0.83	0.83	0.83	6098

```
In [125]: disp = plot_confusion_matrix(baselinexgb, X_test, y_test,  
                                     cmap=plt.cm.Blues,  
                                     xticks_rotation='vertical')  
disp.ax_.set_title('XGB Classifier')  
  
print(disp.confusion_matrix)
```

```
[[1640  147  183]  
 [ 141 1859   54]  
 [ 389  146 1539]]
```



XGBoost is performing better than Random forest and decision Trees in terms of accuracy and precision. Since we are focussing more on lowering the no of false positives XGBoost performs better than KNN.

## Comparing all the model scores:

XGBoost performed the best in terms of it's precision and accuracy.

# Improvements

Since using GridSearchCV take up a lot of computational time, I couldn't put in more parameters like I wanted. Searching out more parameters might improve the model performance.

The model ability to predict non-functional needs repair is still lacking. I could try scrubbing the data in a different way to make the model better.

Try out different types of filling missing values to see if it improves the model.

Instead of dropping the categorical variables like permit, and other duplicate columns we can find better ways to handle those data.

## Next Steps:

We suggest adding using government oversight to make sure wells are installed at lower levels where possible.

Our recommendation is to setup a new fund / organisation to review and repair older wells in cooperation with non-governmental organisations assisting with wells construction.

Payment is a large issues especially when the average monthly salary in Tanzania is \$50.

- Financial support systems from local government areas or districts in addition to payment arrangements in place is crucial to driving a reliable water system.
- Providing more continuous data will be more helpful.

In [ ]: